



南方科技大学

MAT8034: Machine Learning

Classification and Logistic Regression

Fang Kong

<https://fangkongx.github.io/Teaching/MAT8034/Spring2025/index.html>


Outline

- Logistic regression
- Digression: the perceptron learning algorithm
- Newton's method
- Multi-class classification

Logistic regression

Intuition of logistic regression

- Hope to use the linear method to solve the classification problem
- Given a training set: $\{(x^{(i)}, y^{(i)})\}, i = 1, 2, \dots, n$, let $y^{(i)} \in \{0, 1\}$



- Build the connection between p and $\theta^T x = \theta_0 + \theta_1 x_1 + \theta_2 x_2$
- $p \in (0, 1)$ but $\theta^T x \in (-\infty, +\infty)$

Intuition of logistic regression

- Consider the odd: $p/(1-p) \in (0, +\infty)$
- Consider the log odd:
 - $\text{Logit}(p) = \log p/(1-p) \in (-\infty, +\infty)$
- Good properties:
 - $p \rightarrow 0$, $\text{logit} \rightarrow -\infty$; $p \rightarrow 1$, $\text{logit} \rightarrow +\infty$
 - Symmetry: $\text{Logit}(p) = -\text{Logit}(1-p)$
 - Use linear model to approximate the logit: $\theta^\top x = \text{Logit}(p) = \log p/(1-p)$
 - $p = \frac{1}{1 + \exp(-\theta^\top x)} := \text{sigmoid}(\theta^\top x)$


Logistic Regression

- Given a training set $\{(x^{(i)}, y^{(i)}) \text{ for } i = 1, \dots, n\}$ let $y^{(i)} \in \{0, 1\}$.
Want $h_{\theta}(x) \in [0, 1]$. Let's pick a smooth function:

$$h_{\theta}(x) = g(\theta^T x)$$

- Here, g is a link function. There are *many*... but we'll pick one!

$$g(z) = \frac{1}{1 + e^{-z}}.$$



How do we interpret $h_{\theta}(x)$?

$$P(y = 1 \mid x; \theta) = h_{\theta}(x)$$

$$P(y = 0 \mid x; \theta) = 1 - h_{\theta}(x)$$

Likelihood function

- Let's write the Likelihood function. Recall:

$$P(y = 1 \mid x; \theta) = h_{\theta}(x)$$

$$P(y = 0 \mid x; \theta) = 1 - h_{\theta}(x)$$

- Then,

$$\begin{aligned} L(\theta) &= P(y \mid X; \theta) = \prod_{i=1}^n p(y^{(i)} \mid x^{(i)}; \theta) \\ &= \prod_{i=1}^n h_{\theta}(x^{(i)})^{y^{(i)}} (1 - h_{\theta}(x^{(i)}))^{1-y^{(i)}} \quad \text{exponents encode "if-then"} \end{aligned}$$

- Taking logs to compute the log likelihood $\ell(\theta)$ we have:

$$\ell(\theta) = \log L(\theta) = \sum_{i=1}^n y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))$$

Gradient ascent for log likelihood

$$\begin{aligned}\frac{\partial}{\partial \theta_j} \ell(\theta) &= \left(y \frac{1}{g(\theta^T x)} - (1 - y) \frac{1}{1 - g(\theta^T x)} \right) \frac{\partial}{\partial \theta_j} g(\theta^T x) \\ &= \left(y \frac{1}{g(\theta^T x)} - (1 - y) \frac{1}{1 - g(\theta^T x)} \right) g(\theta^T x)(1 - g(\theta^T x)) \frac{\partial}{\partial \theta_j} \theta^T x \\ &= (y(1 - g(\theta^T x)) - (1 - y)g(\theta^T x)) x_j \\ &= (y - h_\theta(x)) x_j\end{aligned}$$

$$\theta_j := \theta_j + \alpha (y^{(i)} - h_\theta(x^{(i)})) x_j^{(i)}$$

Another view: logistic loss

- In linear regression

- The loss function is $J(h_{\theta}(x^{(i)}), y) = (h_{\theta}(x^{(i)}) - y^{(i)})^2$

- For the classification

- Define the loss function

$$\ell_{\text{logistic}}(t, y) \triangleq y \log(1 + \exp(-t)) + (1 - y) \log(1 + \exp(t)). \quad (2.3)$$

- When $y = 1$, minimizing the loss gets $t \rightarrow +\infty, p \rightarrow 1$
- When $y = 0$, minimizing the loss gets $t \rightarrow -\infty, p \rightarrow 0$

Another view: logistic loss

- For the classification
 - Define the loss function

$$\ell_{\text{logistic}}(t, y) \triangleq y \log(1 + \exp(-t)) + (1 - y) \log(1 + \exp(t)) . \quad (2.3)$$

- The relationship between the loss and log likelihood $-\ell(\theta) = \ell_{\text{logistic}}(\theta^\top x, y)$

$$\frac{\partial \ell_{\text{logistic}}(t, y)}{\partial t} = y \frac{-\exp(-t)}{1 + \exp(-t)} + (1 - y) \frac{1}{1 + \exp(-t)} \quad (2.5)$$

$$= 1/(1 + \exp(-t)) - y. \quad (2.6)$$

Then, using the chain rule, we have that

$$\frac{\partial}{\partial \theta_j} \ell(\theta) = -\frac{\partial \ell_{\text{logistic}}(t, y)}{\partial t} \cdot \frac{\partial t}{\partial \theta_j} \quad (2.7)$$


$$= (y - 1/(1 + \exp(-t))) \cdot x_j = (y - h_\theta(x))x_j , \quad (2.8)$$

Connection with the perceptron

Biological neuron structure


■ 细胞结构

- 细胞体
- 树突
- 轴突
- 突触末梢



Biological neural communication

- 细胞膜间的电位表现出的电信号称为动作电位
- 电信号从细胞体中产生，沿着轴突往下传，并且导致突触末梢释放神经递质介质
- 介质通过化学扩散从突触传递到其他神经元的树突
- 神经递质可以是兴奋的或者是抑制的
- 如果从其他神经元来的神经递质是兴奋的且超过某个阈值，将会触发一个动作电位



McCulloch-Pitts neuron model [1943]



- Model the network as a graph, where the units are nodes, and the synaptic connections are weighted edges from node i to node j , with the weight as $w_{j,i}$

- The input of the unit is:


$$\text{net}_j = \sum_i w_{j,i} \cdot o_i$$

- The output of the unit is:

- 0 if $\text{net}_j < T_j$; 1 otherwise
- T_j is the threshold



Single-layer perception by Rosenblatt [1958]




$$\hat{y} = \varphi\left(\sum_{i=1}^m w_i x_i + b\right)$$

$$\varphi(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ -1 & \text{otherwise} \end{cases}$$

- Rosenblatt [1958] 进一步提出感知机作为第一个在“老师”指导下进行学习的模型（即监督学习）
- 专注在如何找到合适的用于二分类任务的权重 w_m
 - $y = 1$: 类别1
 - $y = -1$: 类别2

Training perceptron



预测

激活函数

$$\hat{y} = \varphi\left(\sum_{i=1}^m w_i x_i + b\right)$$

$$\varphi(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ -1 & \text{otherwise} \end{cases}$$

□ 训练

$$w_i = w_i + \eta(y - \hat{y})x_i$$

$$b = b + \eta(y - \hat{y})$$

□ 下列规则等价:

- 如果输出正确, 则不进行操作
- 如果输出高了, 降低正输入的权重
- 如果输出低了, 增加正输入的权重

Newton's method

Another algorithm to maximize $\ell(\theta)$




Newton's method: formulation

- Returning to logistic regression with $g(z)$ being the sigmoid function
- A different algorithm for maximizing the log likelihood $\ell(\theta)$
- To maximize $\ell(\theta)$, hope to find θ such that $\nabla \ell(\theta) = 0$
- New formulation

Given $f : \mathbb{R}^d \rightarrow \mathbb{R}$ find θ s.t. $f(\theta) = 0$.

Newton's method

Given $f : \mathbb{R}^d \rightarrow \mathbb{R}$ find θ s.t. $f(\theta) = 0$



Newton's method

- Suppose $\theta_n - \theta_{n+1} = \Delta$
- $\frac{f(\theta_n) - 0}{\Delta} = f'(\theta_n)$
- $\theta_n - \theta_{n+1} = \Delta = \frac{f(\theta_n)}{f'(\theta_n)}$
- So the update rule in 1d $\theta := \theta - \frac{f(\theta)}{f'(\theta)}$
- To maximizing the log likelihood? $\theta := \theta - \frac{\ell'(\theta)}{\ell''(\theta)}$

Generalization to the multidimensional setting

- For the likelihood, i.e., $f(\theta) = \nabla_{\theta}\ell(\theta)$ we need to generalize to a vector-valued function which has:

$$\theta^{(t+1)} = \theta^{(t)} - \left(H(\theta^{(t)})\right)^{-1} \nabla_{\theta}\ell(\theta^{(t)}).$$

in which $H_{i,j}(\theta) = \frac{\partial}{\partial\theta_i\partial\theta_j}\ell(\theta)$.

Properties of Newton's method

- Convergence rate?
 - Use the Hessian information to determine step size, more adaptive
 - May converge very fast
- Computational cost?
 - Computing Hessian requires $O(d^2)$

Multi-class classification

Problem formulation

- Suppose we want to choose among k discrete values, e.g., $\{'Cat', 'Dog', 'Car', 'Bus'\}$ so $k = 4$.
- We encode with **one-hot** vectors i.e. $y \in \{0, 1\}^k$ and $\sum_{j=1}^k y_j = 1$.

$$\begin{array}{cccc} \begin{pmatrix} 1 \\ 0 \\ 0 \\ 0 \end{pmatrix} & \begin{pmatrix} 0 \\ 1 \\ 0 \\ 0 \end{pmatrix} & \begin{pmatrix} 0 \\ 0 \\ 1 \\ 0 \end{pmatrix} & \begin{pmatrix} 0 \\ 0 \\ 0 \\ 1 \end{pmatrix} \\ 'Cat' & 'Dog' & 'Car' & 'Bus' \end{array}$$

- In this case, $p(y|x; \theta)$ is a distribution over k discrete outcomes

Objective

- Introduce $\theta_1^\top x, \theta_2^\top x, \dots, \theta_k^\top x$ to represent the corresponding probabilities
- Hope:
 - Each probability $\in [0,1]$
 - The sum over all probabilities is 1

Softmax function

- Define the softmax function $\text{softmax} : \mathbb{R}^k \rightarrow \mathbb{R}^k$ as

$$\text{softmax}(t_1, \dots, t_k) = \begin{bmatrix} \frac{\exp(t_1)}{\sum_{j=1}^k \exp(t_j)} \\ \vdots \\ \frac{\exp(t_k)}{\sum_{j=1}^k \exp(t_j)} \end{bmatrix}. \quad (2.9)$$

- Let $(t_1, \dots, t_k) = (\theta_1^\top x, \dots, \theta_k^\top x)$

$$\begin{bmatrix} P(y = 1 \mid x; \theta) \\ \vdots \\ P(y = k \mid x; \theta) \end{bmatrix} = \text{softmax}(t_1, \dots, t_k) = \begin{bmatrix} \frac{\exp(\theta_1^\top x)}{\sum_{j=1}^k \exp(\theta_j^\top x)} \\ \vdots \\ \frac{\exp(\theta_k^\top x)}{\sum_{j=1}^k \exp(\theta_j^\top x)} \end{bmatrix}$$

Quiz

- Does $k = 2$ case agree with logistic regression?

$$P(y = j|x; \theta) = \frac{e^{\theta_j^T x}}{e^{\theta_1^T x} + e^{\theta_2^T x}}$$

How to optimize?

- Compute the negative log likelihood function

$$-\log p(y \mid x, \theta) = -\log \left(\frac{\exp(t_y)}{\sum_{j=1}^k \exp(t_j)} \right) = -\log \left(\frac{\exp(\theta_y^\top x)}{\sum_{j=1}^k \exp(\theta_j^\top x)} \right)$$

- Define the cross-entropy loss function

$$\ell_{\text{ce}}((t_1, \dots, t_k), y) = -\log \left(\frac{\exp(t_y)}{\sum_{j=1}^k \exp(t_j)} \right)$$

- Over n training examples?

$$\ell(\theta) = \sum_{i=1}^n \ell_{\text{ce}}((\theta_1^\top x^{(i)}, \dots, \theta_k^\top x^{(i)}), y^{(i)})$$

Gradient descent to minimize the loss

$$\frac{\partial \ell_{\text{ce}}(t, y)}{\partial t_i} = \phi_i - 1\{y = i\}, \quad (2.16)$$

where $1\{\cdot\}$ is the indicator function, that is, $1\{y = i\} = 1$ if $y = i$, and $1\{y = i\} = 0$ if $y \neq i$. Alternatively, in vectorized notations, we have the following form which will be useful for Chapter 7:

$$\frac{\partial \ell_{\text{ce}}(t, y)}{\partial t} = \phi - e_y, \quad (2.17)$$

where $e_s \in \mathbb{R}^k$ is the s -th natural basis vector (where the s -th entry is 1 and all other entries are zeros.) Using Chain rule, we have that

$$\frac{\partial \ell_{\text{ce}}((\theta_1^\top x, \dots, \theta_k^\top x), y)}{\partial \theta_i} = \frac{\partial \ell(t, y)}{\partial t_i} \cdot \frac{\partial t_i}{\partial \theta_i} = (\phi_i - 1\{y = i\}) \cdot x. \quad (2.18)$$

Therefore, the gradient of the loss with respect to the part of parameter θ_i is

$$\frac{\partial \ell(\theta)}{\partial \theta_i} = \sum_{j=1}^n (\phi_i^{(j)} - 1\{y^{(j)} = i\}) \cdot x^{(j)}, \quad (2.19)$$

Summary

- Two-class classification
 - Logistic regression
 - Intuition, optimization
 - Digression: the perceptron learning algorithm
 - Newton's method
 - Use second-order information
- Multi-class classification
 - Softmax function